

S3F Barrier Types & Performance

Barrier Types

- Original s3f barrier (barrier.cc)
- Mutex (broadcast) barrier (barrier_mutex.cc)
- Fast sense-reversing barrier (fast_barrier.cc)
 - Fast-userspace locking (futex) option
- Software combining tree (fast_tree_barrier.cc)
 - Arrival futex option
 - Release futex option
- Pthread barrier

Original S3F Barrier

- Pros:
 - No locking at all
- Cons:
 - Freezes after 2^{31} calls on systems with 32-bit longs (or 2^{65} calls for 64-bit longs)
 - Potential invalid behavior on non-TSO (total-store-ordering) architectures (ie non-x86)
 - Extra unnecessary work by all but one thread
 - Poor cache performance due to memory contention

Broadcast Barrier

- Uses `pthread_cond_wait` and `pthread_cond_broadcast` with a mutex lock
- Pros:
 - Sense-reversing (no call-limit)
 - Better cache behavior
- Cons:
 - VERY slow wakeup hinders s3f's performance

'Fast' Sense-Reversing Barrier

- Can use `sched_yield()` OR futex locking
 - Enable `#define USE_FUTEX` in `fast_barrier.h`
- Pros:
 - Sense-reversing (no call limit)
 - Optimized cache performance
 - Only locks to update sim window statistics
- Cons:
 - Can lead to high 'sys' (kernel-space) times (still better than original barrier) for both yield and futex

Software Combining Tree

- Configurable with defaults:
 - `ARRIVAL_CHILDREN = 4`
 - `RELEASE_CHILDREN = 2`
- Can use `sched_yield()` or futex locking (defines in `fast_tree_barrier.h`):
 - For arriving children (`#define USE_ARRIVAL_FUTEX`) – not recommended, generally increases 'sys' time overhead
 - For released children (`#define USE_RELEASE_FUTEX`) – generally safe, but better if `RELEASE_CHILDREN` increased
- Pros:
 - Excellent for high numbers of threads
 - Potentially scalable to a distributed system
- Cons:
 - Extra lock overhead for low numbers of threads
 - Slightly slower release times
 - Futex locking does not appear to perform well

Pthread Barrier

- Uses futex locking via assembly code
- Pros:
 - Ease of use – default pthread implementation
 - Excellent for large amounts of thread work
 - Very low (comparative) 'sys' time
- Cons:
 - Slightly slower wakeup than yield
 - Barrier not guaranteed in all pthread implementations
 - Manual window statistics updating required

Using Alternative Barrier Types

- Enable exactly ONE in s3f.h:
 - '#define MUTEX_BARRIER' – mutex broadcast
 - '#define PTHREAD_BARRIER' – pthread
 - (0) '#define SCHED_YIELD_BARRIER' – all others
- Additional toggles:
 - (1) '#define FAST_SCHED_YIELD_BARRIER', requires (0)
 - (2) '#define TREE_BARRIER', requires (0) and (1)
 - (0) only → original barrier
 - (0) + (1) only → 'fast' sense-reversing barrier
 - (0) + (1) + (2) → software combining tree barrier
- Enable/disable futexes within *barrier* headers

Performance Summary

- Pthread barrier adds the least overhead for almost all cases
- fast_barrier works well for very low thread work
- sched_yield() works best for fewer threads
- Futex locking works better for very small simulation synchronization windows (SSW)
 - WARNING: 'sys' time scales up exponentially as SSW increases
- fast_tree_barrier looks to work well for 32+ threads
 - Should be more extensible to a distributed system
- All performance looks to scale linearly with number of hosts
- Detailed performance data available in barrier_performance.xls (in barriers.tar.gz)